


Published baselines against our single-disease encoder, cell-pooled RMSE



Every model in this chart trains on the disease it is scored on. Bars are the mean over 5 seeds, normalised per row so our single-disease encoder is 1.0; longer is worse.

* dengue is pooled over the 2,392-node subsample on BOTH sides, because the baseline repos cannot take 7,165 nodes. EpiGNN is the only baseline that runs on dengue, so the Cola-GNN and HeatGNN bars are absent there by construction, not by omission. MTGNN is excluded as a comparator: 47 of its 80 prediction files hold one repeated constant. COVID has no baseline of any kind and is not shown; Ebola has none either, since no node-indexed baseline can run on a disease with no training data.