





A · share of IG attribution by input channel



B · share of IG attribution by lag band, per panel

C · signed IG on Montserrado, Liberia's OWN h3 forecast
target week 2014-10-25: forecast 31.4 cases/week against 1,428 observed

D · neighbour influence on Montserrado, Liberia, edge ablation



Integrated gradients from a zero baseline over the [N, 20, 4] trunk input, from the saved checkpoints: 5 development panels and 2 Ebola L12 arms, 5 seeds. A and B are shares of absolute IG attribution, pooled over scored origins and horizons.

B is reported at lag BAND level: the 20-lag profile combs with period 4 on every panel and untrained encoders reproduce the comb, so single-lag resolution reads the dilated TCN's receptive field, not epidemiology.

obs_mask is constant on the three influenza panels and COVID, so nothing there varies it. Its share is still large (up to 0.18) because a zero baseline charges the channel for the whole 0 -> 1 move: an artefact of the baseline, not a use of observation status.

C and D are one district at the national peak week. C targets that district's OWN median forecast, not the national sum: its h3 forecast is 31.4 cases/week against 1,428 observed for 2014-10-25.

That week carries a whole multi-week reporting increment (gap-lumping, client_decisions.md A4), so the observed peak is inflated and the quiet weeks either side are flattened.

A bar at exactly zero is a neighbour that filed no report that week: each edge is weighted by the neighbour's own observation mask, so its message is already zero and dropping it changes nothing.

Edge ablation holds every degree fixed, so only the neighbour message moves. None of D is a source of accuracy: the gate-off ablation shows neighbour information helps error in 0 of 40 cells.

That ablation zeroes the gate but keeps the LTR degree feature, so the tally bounds neighbour information, not the graph in total. D shows where the model draws from, and nothing more.