


What holding the disease out costs, cell-pooled RMSE
hatched bars never see the target disease, the other two train on it



The two hatched bars are a DIFFERENT task, not a fairer version of the same one. Blue and black train on the disease they are scored on; orange holds that disease out of training entirely. EpiGNN, Cola-GNN and HeatGNN are inductive and node-indexed, so they cannot be run in the held-out setting at all and have no bar there. Read the orange bars as the price of the setting, not as a race. Baselines are collapsed to the best one per cell, which steel-mans them. Bars are the mean over 5 seeds, normalised per row so our single-disease encoder is 1.0. * dengue is pooled over the 2,392-node subsample on all bars. The LDO3 records predate the trunk early-stop fix; a full-budget rerun at seed 42 reproduced them bit-identically on all three folds (10 of 10), but that check is seed 42 only. COVID is not shown because it has no baseline; Ebola is not shown because no node-indexed baseline can run on it.